

## 实验二、 MapReduce 的实现

### 一、流程及代码分析：

要求：通过流程图的形式，结合代码实现，分析 MapReduce 执行过程，并给出结果截图。

本次实验中的部分代码细节与原论文中的模型有所出入，例如：论文中的模型为 master 主动 ping 各个 worker 并分配任务，而在本次实验中采取的做法是由各个 worker 主动向 master 发起请求，master 将分配的任务信息放在 reply 中回复给 worker，这样做的好处是实现起来更为简便，只有 master 需要启动一个 rpc 服务进行监听，而 worker 作为 rpc 客户端调用即可。

#### master 流程分析：

master（代码中叫 coordinator，二者是一个东西）只负责分配任务，由于 master 中并行执行了多个线程，就不对 master 画流程图了，直接用文字结合代码说明，后面的 worker 再画图：

1. 启动 master 并初始化参数，启动超时检查线程和 rpc 监听线程；

```
// MakeCoordinator : create a Coordinator.  
// main/mrcoordinator.go calls this function.  
// nReduce is the number of reduce tasks to use.  
func MakeCoordinator(files []string, nReduce int) *Coordinator { 2 usages & lvyy1999  
    // init master params  
    c := Coordinator{...}  
  
    // init map tasks  
    for i, file := range files {...}  
  
    // init reduce tasks  
    for i := 0; i < nReduce; i++ {...}  
  
    // run a thread to check task timeout  
    go c.CheckTaskTimeout()  
    // run a rpc server to listen request from worker  
    c.server()  
    return &c  
}
```

2. 定期检查是否有超时任务，若有，将其状态重置为空闲以便重新分配；

```
// CheckTaskTimeout : check if any task not completed in specified time  
// if a task is timeout, reset its state, make it available for another worker  
func (c *Coordinator) CheckTaskTimeout() { 2 usages & lvyy1999  
    for !c.Done() {  
        // get current time  
        current := time.Now().Unix()  
  
        c.mMutex.Lock()  
        // check map tasks  
        for idx := range c.mMapTasks {...}  
        // check reduce tasks  
        for idx := range c.mReduceTasks {...}  
        c.mMutex.Unlock()  
  
        // wait a little time to check again  
        time.Sleep(time.Second * time.Duration(CheckInterval))  
    }  
}
```

3. 当有 worker 来申请任务时，分配一个可执行的任务并更新任务状态。若有空闲的 map 任务则先分配 map 任务，只有当所有 map 任务都执行完才可以分配 reduce 任务，对于 reduce 任务，需要将 map 任务产生的中间文件地址告诉执行 reduce 任务的 worker；

```
func (c *Coordinator) AskForTaskHandler(args *AskForTaskArgs, reply *AskForTaskReply) error { 1 usage  new *
    c.mMutex.Lock()
    defer c.mMutex.Unlock()
    // if the workerID is invalid, assign a new workerID
    workerID := args.WorkerID
    if workerID == 0 {...}
    reply.WorkerID = workerID

    // choose an available task to assign
    var task *TaskInfo
    if c.mNUncompletedMap > 0 { // if there are some map tasks uncompleted, can only assign map task
        for i := range c.mMapTasks {...}
    } else if c.mNUncompletedReduce > 0 { // after all map tasks completed, can assign reduce task
        for i := range c.mReduceTasks {...}
    }

    if task == nil {...} else {
        // pass the available task info to reply
        reply.TaskID = task.TaskID
        reply.NReduce = c.mNReduce
        reply.TaskType = task.TaskType
        reply.InputFile = task.InputFile
        reply.IntermediateFiles = task.IntermediateFiles
        // update task state
        task.WorkerID = workerID
        task.StartTime = time.Now().Unix()
        task.TaskState = TaskStateInProgress
    }
}
```

4. 当有 worker 来提交任务时，检查信息是否匹配（如提交者的 id 和本地记录的执行者的 id 是否一致等），然后根据任务类型做相应处理（如若为 map 任务，需记录中间文件地址）；

```
// SubmitTaskHandler : handle the SubmitTask request from worker
func (c *Coordinator) SubmitTaskHandler(args *SubmitTaskArgs, reply *SubmitTaskReply) error { 1 usage  &lvyy1999 *
    c.mMutex.Lock()
    defer c.mMutex.Unlock()
    // choose the tasks list by args.TaskType
    var tasks []*TaskInfo
    if args.TaskType == TaskTypeMap {
        tasks = &c.mMapTasks
    } else if args.TaskType == TaskTypeReduce {
        tasks = &c.mReduceTasks
    } else {
        fmt.Println("... unknown task type")
        return nil
    }

    // check the taskID, and choose the corresponding task
    var task *TaskInfo
    taskID := args.TaskID
    if taskID < 0 || taskID >= len(*tasks) {
        fmt.Println("... task id out of range")
        return nil
    } else {
        task = &(*tasks)[taskID]
    }
```

```

// check the workerID with recent allocated workerID
// because the task maybe reallocated to another worker when timeout
if args.WorkerID != (*task).WorkerID {
    fmt.Println("... workerID not match")
    return nil
}

// check the task state, can only transfer to completed from in progress
if (*task).TaskState != TaskStateInProgress {
    fmt.Println("... task state not match")
    return nil
}

// pass all check, update task and master state, return result to worker
reply.TaskState = TaskStateCompleted
(*task).TaskState = TaskStateCompleted
if args.TaskType == TaskTypeMap {
    c.mUncompletedMap--
    // if map completed, store intermediate files to reduce task info
    for reduceTaskID, file := range args.IntermediateFiles {
        c.mReduceTasks[reduceTaskID].IntermediateFiles = append(c.mReduceTasks[reduceTaskID].IntermediateFiles, file)
    }
} else {
    c.mUncompletedReduce--
}

return nil

```

5. 当所有任务都执行完时，就可以退出了（这里不是直接退出，是模型调用者会不断调用 master 的 Done 接口查询工作完成情况，当返回 true 时调用者会退出），由于 map 任务都完成后才会分配 reduce 任务，所以这里判断 reduce 任务有没有全部完成就行了；

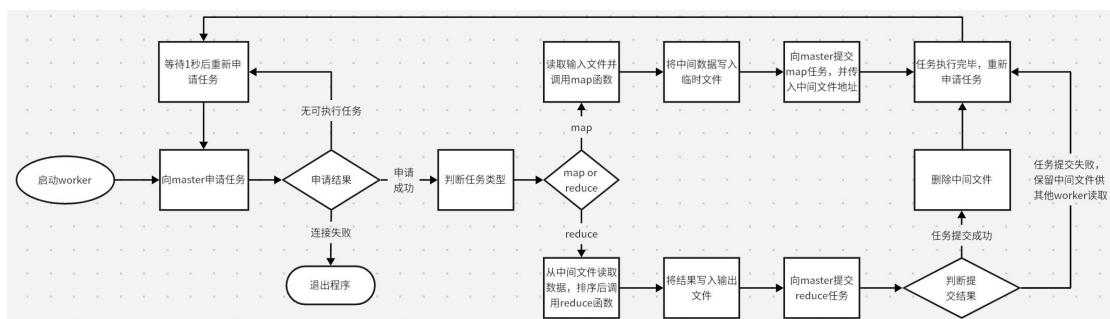
```

// Done : return true when all tasks completed
// main/mrcoordinator.go calls Done() periodically to find out
// if the entire job has finished.
func (c *Coordinator) Done() bool { 4 usages & lvyy1999
    c.mMutex.Lock()
    ret := c.mUncompletedReduce == 0
    c.mMutex.Unlock()
    return ret
}

```

## worker 流程分析：

worker 的流程相较 master 要复杂一些，主要流程图如下：



## 1. 启动 worker 并不断向 master 申请任务，根据任务类型转入相应处理函数；

```
for {
    // call rpc method to ask for a task
    args := AskForTaskArgs{}
    args.WorkerID = workerID
    reply := AskForTaskReply{}
    ok := call(rpcname: "Coordinator.AskForTaskHandler", &args, &reply)
    if !ok {...} else {
        // call successfully, restart counting failure times
        failedCount = 0

        // be assigned a workerID by master
        if workerID == 0 {...}

        // do task according to reply.TaskType
        switch reply.TaskType {
        case TaskTypeNone:
            // no task to do
        case TaskTypeMap:
            DoMapTask(mp, reply.InputFile, reply.TaskID, reply.WorkerID, reply.NReduce)
        case TaskTypeReduce:
            DoReduceTask(reducef, reply.IntermediateFiles, reply.TaskID, reply.WorkerID, reply.NReduce)
        }
    }

    // wait a little time to ask again
    time.Sleep(time.Second * time.Duration(AskForTaskInterval))
}
```

## 2. 对于 map 任务：

### 2.1 读取输入文件并调用 client 提供的 map 方法；

```
// read file and pass it to Map
file, err := os.Open(filename)
if err != nil { log.Fatalf("cannot open %v", filename) }
content, err := ioutil.ReadAll(file)
if err != nil { log.Fatalf("cannot read %v", filename) }
err = file.Close()
if err != nil { fmt.Println("cannot close file, err:", err) }
// call the map func from the client
kva := mpf(filename, string(content))
```

### 2.2 将中间数据根据 key 进行 hash 分区，并以 json 格式写入相应的中间文件；

```
// write intermediate data into files from memory
intermediateFiles := make(map[int]string)
for i := 0; i < nReduce; i++ {
    // mr-X-Y, where X is the Map task number, and Y is the reduce task number
    outfileName := "mr-" + strconv.Itoa(taskID) + strconv.Itoa(i)
    // create temp intermediate file to store intermediate data
    outfile, err := ioutil.TempFile("", outfileName)
    if err != nil { log.Fatalf("cannot create temp %v", outfileName) }
    // defer to close the temp file
    // can't delete it in this task, because the reduce worker need to read it later
    defer outfile.Close()
    // record th intermediate files' name, and return them to master later
    intermediateFiles[i] = outfile.Name()
    // encode json and write into corresponding intermediate file
    enc := json.NewEncoder(outfile)
    for _, kv := range kva {
        // use ihash(key) % NReduce to choose the reduce task number
        if ihash(kv.Key)%nReduce == i {
            if err := enc.Encode(&kv); err != nil {
                fmt.Println("encode json error:", err)
            }
        }
    }
}
```

2.3 通知 master 任务执行完毕，并传入中间文件的路径地址；

```
// notify the master that this map task has completed
args := SubmitTaskArgs{
    TaskID:           taskID,
    WorkerID:         workerID,
    TaskType:         TaskTypeMap,
    IntermediateFiles: intermediateFiles,
}
reply := SubmitTaskReply{}
ok := call(rpcname: "Coordinator.SubmitTaskHandler", &args, &reply)
if !ok {
    fmt.Println(a... "call Coordinator.SubmitTaskHandler failed!")
}
```

3. 对于 reduce 任务：

3.1 从 master 传入的中间文件地址读取 json 并解析成键值对；

```
// read each intermediate files, read intermediate data from the json
intermediate := []KeyValue{}
for _, filename := range intermediateFiles {
    // open the intermediate file
    file, err := os.Open(filename)
    if err != nil { log.Fatalf(format: "cannot open %v", filename) }
    defer file.Close()

    // read json from file and decode
    dec := json.NewDecoder(file)
    for {
        var kv KeyValue
        if err := dec.Decode(&kv); err != nil { break }
        intermediate = append(intermediate, kv)
    }
}
```

3.2 对中间数据（键值对）进行排序；

```
// sort the intermediate key-value pairs by key
sort.Sort(ByKey(intermediate))
```

3.3 先创建一个临时文件（防止文件在全部写入完成之前就被测试程序看到）；

```
// To ensure that nobody observes partially written files in the presence of crashes,
// create a temporary file firstly and atomically rename it when it is completely written.
outfileName := "mr-out-" + strconv.Itoa(taskID)
outfile, err := ioutil.TempFile(dir: "", outfileName)
if err != nil {
    log.Fatalf(format: "cannot create temp file %v", outfileName)
}
defer outfile.Close()
```

3.4 将属于同一个 key 的 values 传入 client 提供的 reduce 函数，并将结果写入临时文件；

```
// call Reduce on each distinct key in intermediate[],
// and print the result to mr-out-Y.
for i, j := 0, 1; i < len(intermediate); i, j = j, j+1 {
    for j < len(intermediate) && intermediate[j].Key == intermediate[i].Key {
        j++
    }
    var values []string
    for k := i; k < j; k++ {
        values = append(values, intermediate[k].Value)
    }

    // call the reduce func from client
    output := reducef(intermediate[i].Key, values)

    // this is the correct format for each line of Reduce output.
    _, err = fmt.Fprintf(outfile, "%v %v\n", intermediate[i].Key, output)
    if err != nil { log.Fatalf(format: "cannot write to file %v", fileName) }
}
```

3.5 文件写好后，再通过 os.Rename 对其进行原子性地改名，使得其对测试脚本可见；

```
// rename the temp file
err = os.Rename(outfile.Name(), fileName)
if err != nil { log.Fatalf(format: "cannot rename temp file %v", fileName) }
```

3.6 向 master 提交 reduce 任务已完成；

```
// notify the master that this reduce task has completed
args := SubmitTaskArgs{
    TaskID: taskID,
    WorkerID: workerID,
    TaskType: TaskTypeReduce,
}
reply := SubmitTaskReply{}
ok := call(rpcname: "Coordinator.SubmitTaskHandler", &args, &reply)
```

3.7 若任务提交成功（调用 master 的任务提交接口成功且返回的任务状态为已完成），则可以删除 map 任务产生的那些存储中间数据的临时文件，以避免对存储的浪费；

```
if !ok {
    fmt.Println(a... "call Coordinator.SubmitTaskHandler failed!")
} else if reply.TaskState != TaskStateCompleted {
    fmt.Println(a... "not submit successfully, maybe timeout!")
} else {
    // can only clean the intermediate files after completing the reduce task and submit to master successfully,
    // because the master may reassign the task to another worker if this worker is timeout, so that these files may be read again.
    for _, filename := range intermediateFiles {
        err = os.Remove(filename)
        if err != nil { fmt.Println(a... "remove temp file error:", err) }
    }
}
```

4. 若连接 master 失败，则可认为 master 已退出，那么 worker 也可以退出。

```
ok := call(rpcname: "Coordinator.AskForTaskHandler", &args, &reply)
if !ok {
    fmt.Println(a... "call Coordinator.AskForTaskHandler failed!")
    // if call failed three times, maybe master has finished, worker can exit
    if failedCount++; failedCount >= AskForTaskFailedLimit {
        fmt.Println(a... "program will exit!")
        return
    }
}
```

## 二、实验结果截图：

本实验运行环境为 centOS7.4 虚拟机

```
(base) [root@localhost main]# ./test-mr.sh
*** Starting wc test.
--- wc test: PASS
2025/11/03 19:42:35 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:42:36 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:42:36 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
*** Starting indexer test.
2025/11/03 19:42:53 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- indexer test: PASS
2025/11/03 19:42:54 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
*** Starting map parallelism test.
2025/11/03 19:43:12 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:43:13 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- map parallelism test: PASS
*** Starting reduce parallelism test.
2025/11/03 19:43:32 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:43:33 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- reduce parallelism test: PASS
*** Starting job count test.
2025/11/03 19:44:02 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:44:03 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:44:03 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:44:03 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- job count test: PASS
*** Starting early exit test.
2025/11/03 19:44:16 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:44:16 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:44:17 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- early exit test: PASS
*** Starting crash test.
map task 0 timeout, workerID = 1
map task 1 timeout, workerID = 2
map task 2 timeout, workerID = 3
map task 4 timeout, workerID = 5
map task 0 timeout, workerID = 8
map task 7 timeout, workerID = 4
map task 4 timeout, workerID = 6
map task 7 timeout, workerID = 7
map task 4 timeout, workerID = 11
reduce task 5 timeout, workerID = 10
2025/11/03 19:45:12 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:45:12 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
2025/11/03 19:45:13 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused
--- crash test: PASS
*** PASSED ALL TESTS
(base) [root@localhost main]#
```

注：

1. 图片中的 dialing:dial unix /var/tmp/824-mr-0: connect: connection refused 报错是由于 worker 在连接 master 的 rpc 服务失败时才会退出，如果不显示这个报错，可以修改 worker.go 中的 call 函数，去掉这个打印；或者由 master 在任务全部完成后，反向通知各个 worker 主动退出。
2. 项目提供的测试脚本中的 early exit test 一节，用到了 wait -n 命令，当系统的 bash 版本比较老时，会因为不支持这一命令导致这一项测试失败（没有等待任何进程退出就保存了结果，导致报错 output changed after first worker exited），需要将 bash 版本升级到 4.3 以上进行解决。

## 三、完整代码

由于代码较长，这里不贴出，可以直接看文件夹内的源代码文件，或从我的 github 仓库获取：<https://github.com/lvyy1999/My6.824>。